

Curso de

# Redes Neuronales con TensorFlow

Adonaí Vera  
Co-Founder y CTO en SwitchAI



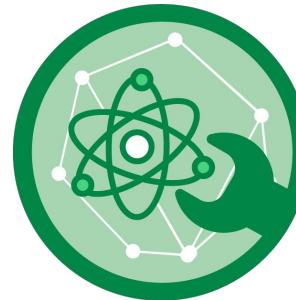
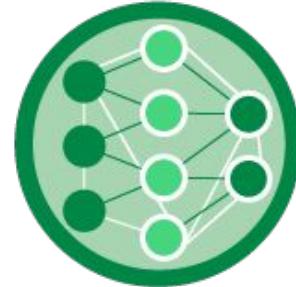
A portrait of a young man with short dark hair and round glasses, smiling broadly. He is wearing a long-sleeved, button-down shirt in a muted olive-green color. The background is plain white.

[www.switchai.co](http://www.switchai.co)



# Requisitos del curso

- Fundamentos de redes neuronales.
- Uso de Jupyter Notebooks y Google Colab.
- Creación de proyectos de ciencia de datos e inteligencia artificial.





# Por qué debes tomar el curso

- Tienes un proyecto de la universidad en donde aplicarás redes neuronales.
- En tu empresa están trabajando con redes neuronales y te asignaron el nuevo proyecto.
- Curioso por nuevas tecnologías.

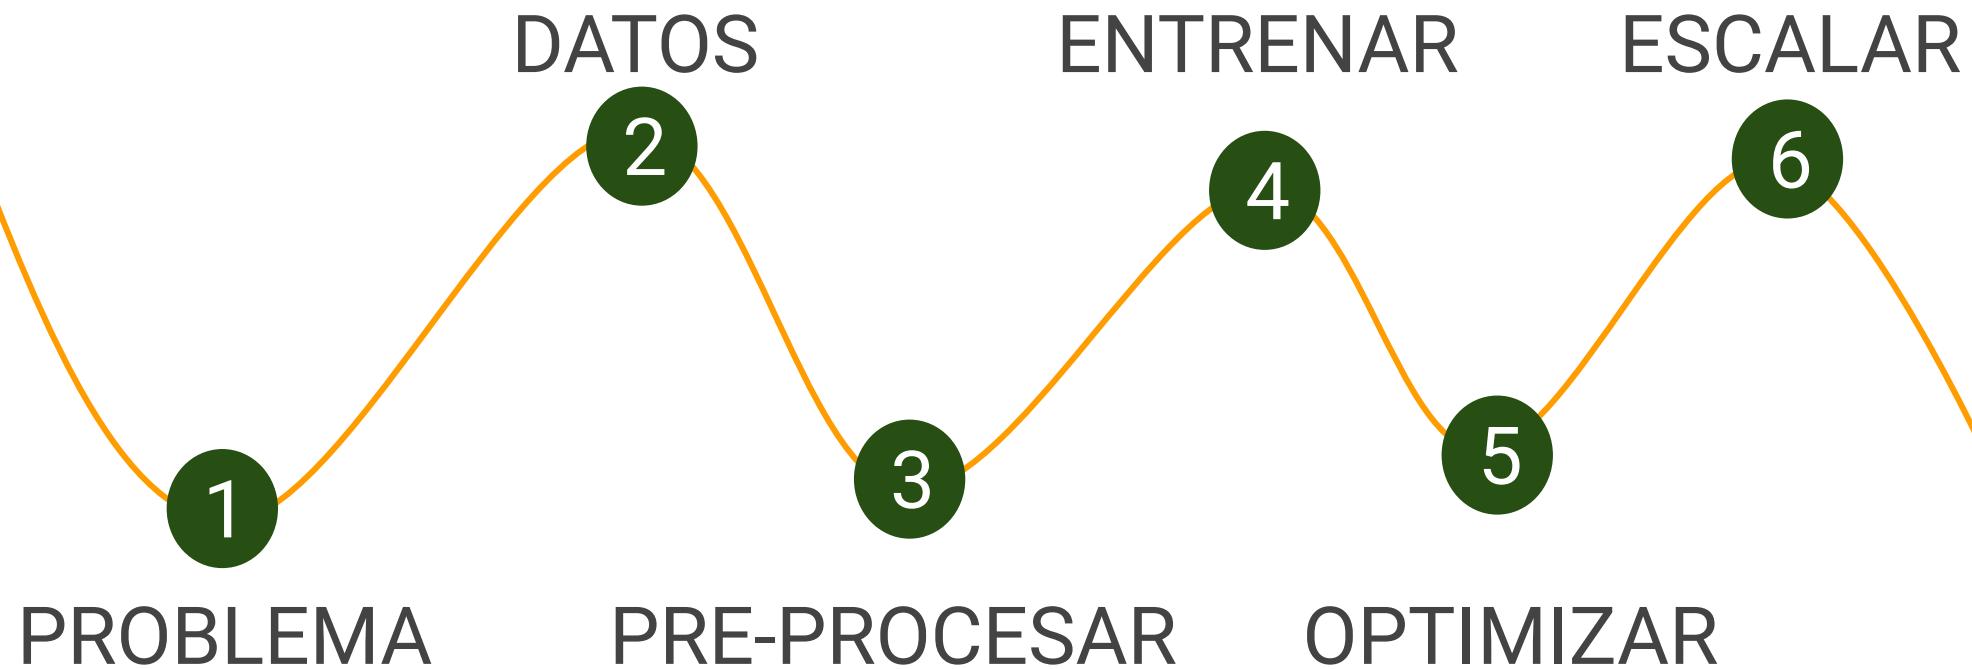


# Por qué debes tomar el curso

- Quieres aprender cómo aplicar inteligencia artificial a imágenes.
- Tienes las bases pero quieres volverte un pro configurando redes neuronales.
- Quieres crear tu propio modelo.



# Ciclo de la IA





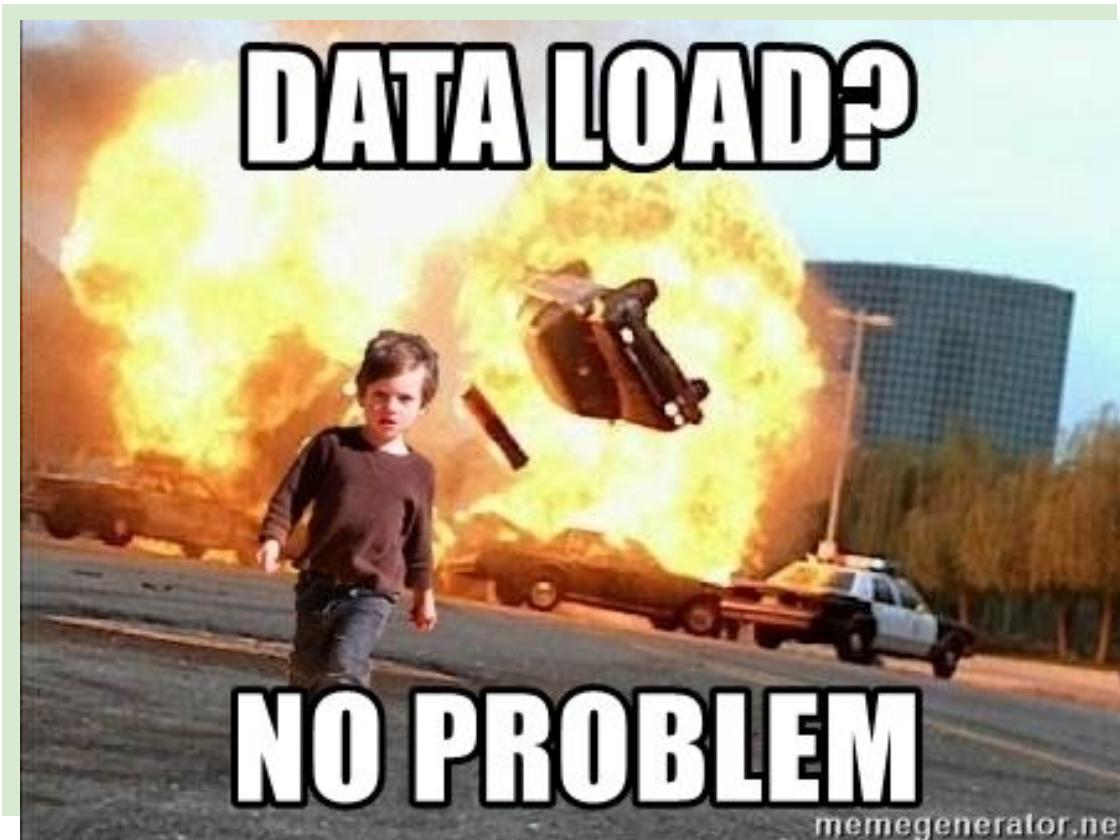
# Redes neuronales

Llevar redes  
neuronales a la  
práctica.



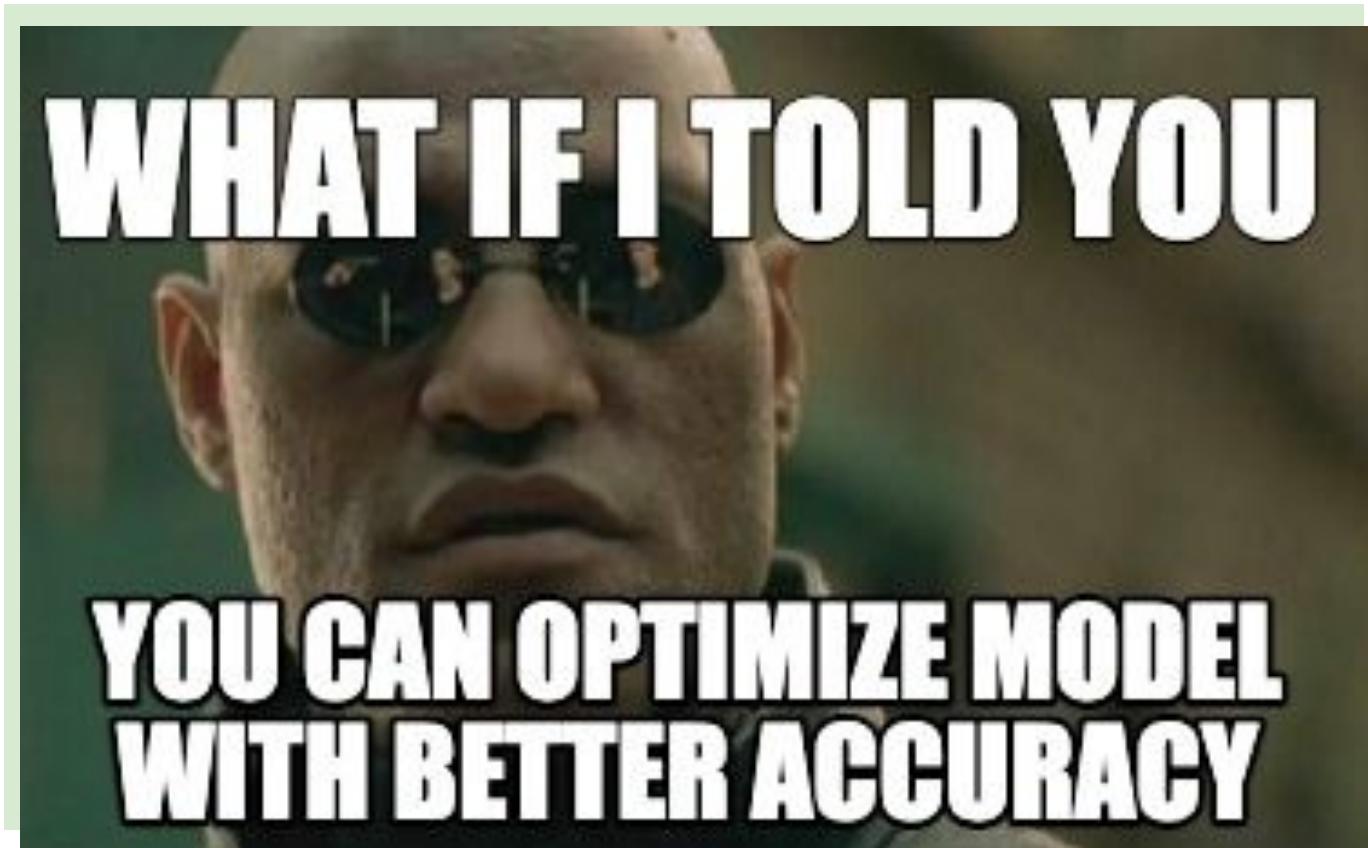


# Cargar bases de datos





# Optimizar modelos





# ¿Cómo evitar el overfitting y underfitting?





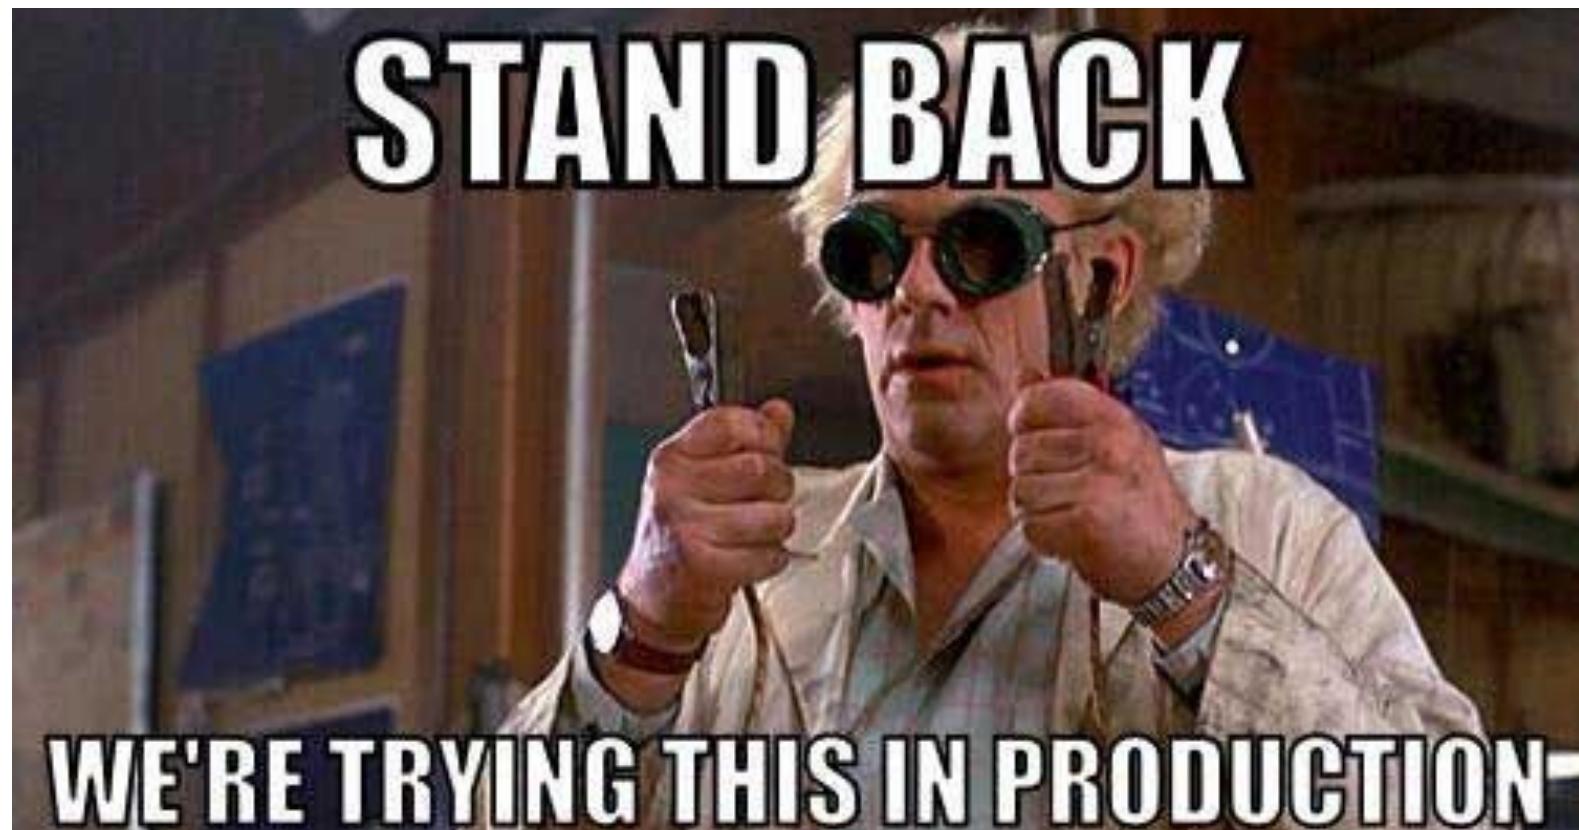
# Transferencia de aprendizaje

Transfer learning be like





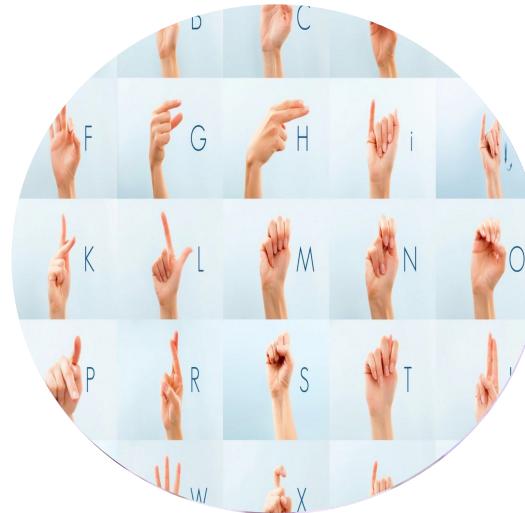
# Almacenar y cargar modelos





# Proyecto del curso

- 27455 imágenes
- Escala de grises
- 28 x 28
- 24 clases (J - Z)
- JPEG
- TecPerson - Kaggle
- C0: Public Domain



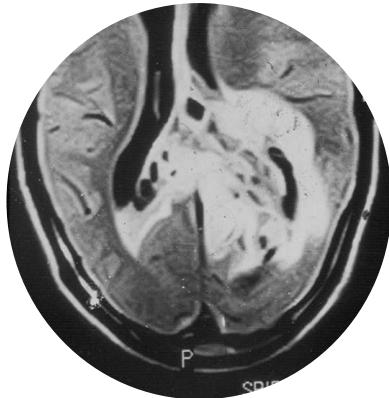
**Lenguaje de señas**



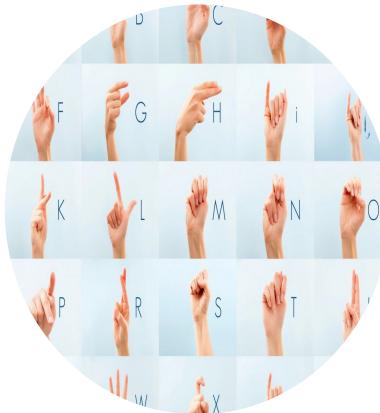
# Proyectos



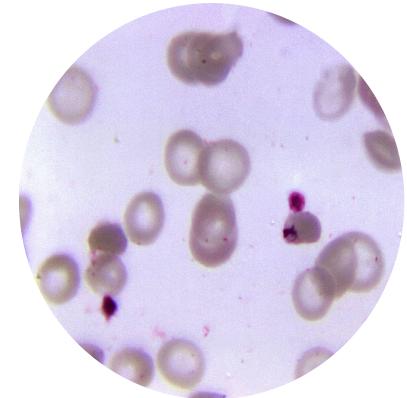
Tom &  
Jerry



Tumor



Lenguaje  
de Señas



Malaria



```
# Bases de datos Malaria
!wget --no-check-certificate \
    https://storage.googleapis.com/platzi-tf2/MalariaCells.zip \
    -O /tmp/sign-language-img.zip

# Bases de datos Tumor
!wget --no-check-certificate \
    https://storage.googleapis.com/platzi-tf2/TumorClassification.zip \
    -O /tmp/sign-language-img.zip

# Bases de datos Tom y Jerry
!wget --no-check-certificate \
    https://storage.googleapis.com/platzi-tf2/TomAndJerry.zip \
    -O /tmp/sign-language-img.zip
```



# ¿Qué aprenderás?

- Cómo cargar tus propias bases de datos.
- Cargar bases de datos en formatos como CSV, JSON, BASE64, imágenes.
- Aplicar técnicas para optimizar tus modelos.



# ¿Qué aprenderás?

- Agregar métricas en el entrenamiento de tus modelos.
- A cargar y guardar modelos.
- Autotunner de Keras para encontrar las mejores variables.



# ¿Qué aprenderás?

- Bases de aprendizaje por transferencia.
- Uso de TensorBoard y cómo mostrar tu proyecto al mundo entero.
- A tener tu modelo listo para utilizarlo como inferencia.



# Introducción a TF2.0



# Librerías de deep learning

Caffe



Microsoft

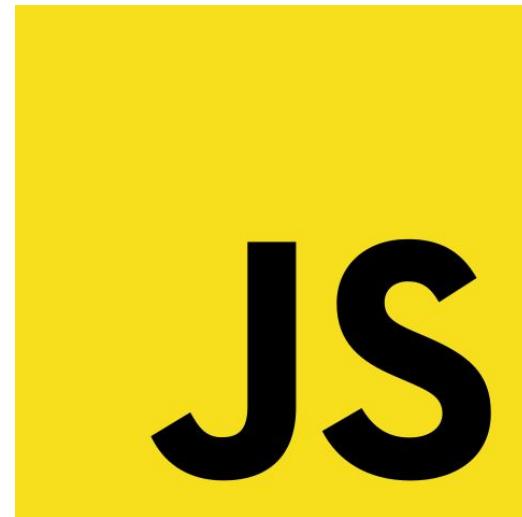
CNTK



GLUON



# Lenguajes de programación



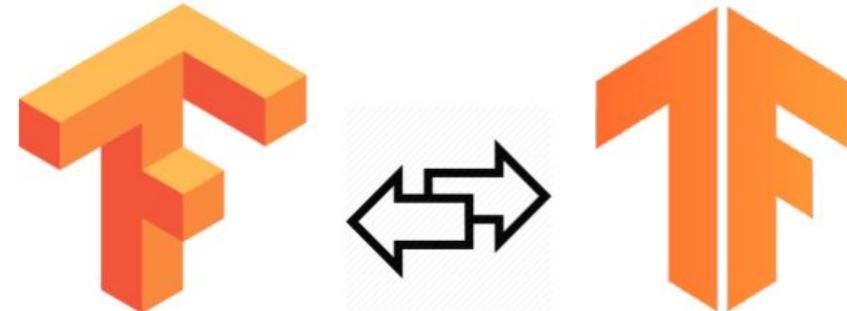


# TensorFlow 2 vs. TensorFlow 1



Eager Execution

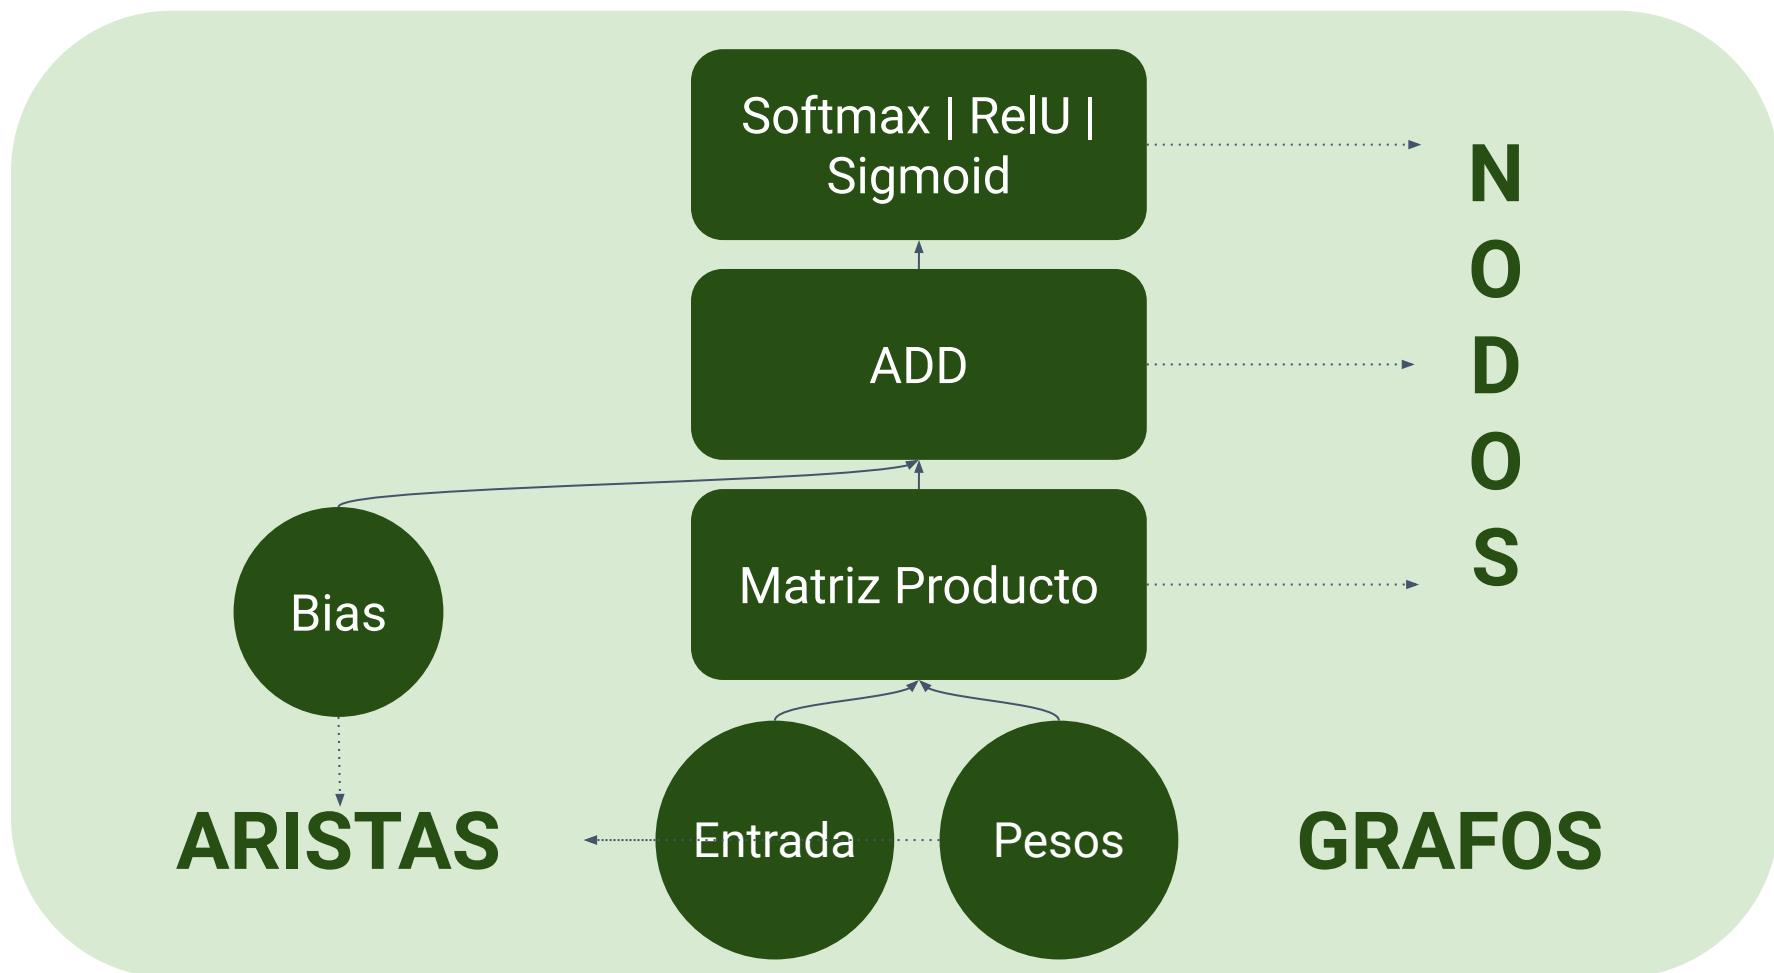
Symbolic | Imperative API



- API CLEAN UP
- NO MORE GLOBAL FUNCTIONS
- NOT SESSIONS



# Lógica de TensorFlow





# Diferencias en código

```
● ● ●

import tensorflow as tf
import tensorflow.compat.v1 as v1
import tensorflow_datasets as tfds

g = v1.Graph()

with g.as_default():
    in_a = v1.placeholder(dtype=v1.float32, shape=(2))
    in_b = v1.placeholder(dtype=v1.float32, shape=(2))

    def forward(x):
        with v1.variable_scope("matmul", reuse=v1.AUTO_REUSE):
            W = v1.get_variable("W", initializer=v1.ones(shape=(2,2)),
                                regularizer=lambda x:tf.reduce_mean(x**2))
            b = v1.get_variable("b", initializer=v1.zeros(shape=(2)))
        return W * x + b

    out_a = forward(in_a)
    out_b = forward(in_b)
    reg_loss=v1.losses.get_regularization_loss(scope="matmul")

with v1.Session(graph=g) as sess:
    sess.run(v1.global_variables_initializer())
    outs = sess.run([out_a, out_b, reg_loss],
                  feed_dict={in_a: [1, 0], in_b: [0, 1]})
```

```
● ● ●

import tensorflow as tf
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test, verbose=2)
```



# TensorFlow + Keras

Código abierto

Google

2015

Escrita en Python

API alto nivel

François  
Chollet

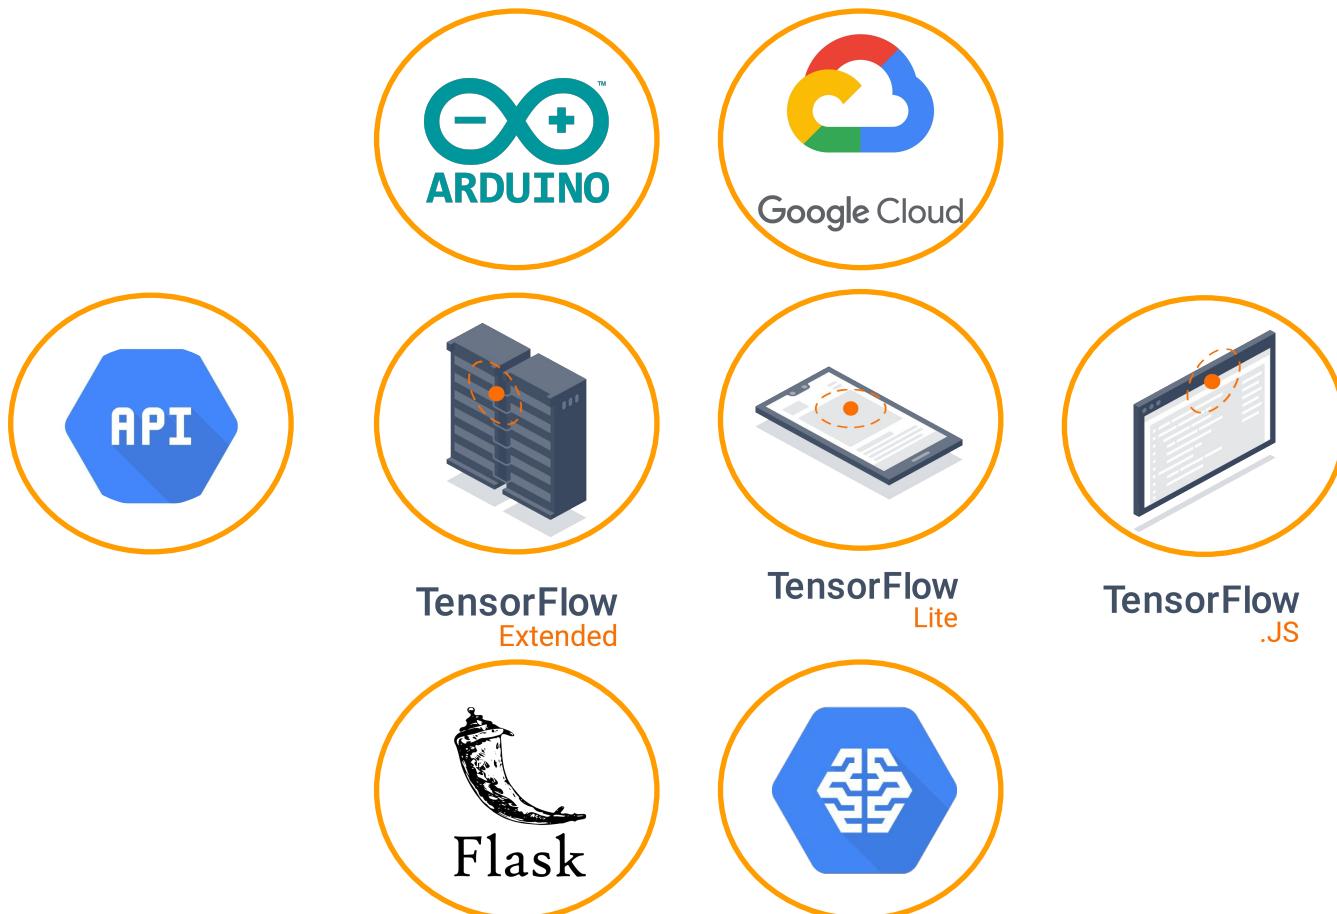
Artificial  
Intelligence

with Lex Fridman





# El ecosistema de TensorFlow





# Comunidad de TF

- **56,000** preguntas en StackOverflow.
- **109** ML Google Developer Experts.
- **46** TensorFlow user groups en todo el mundo.
- **25** publicaciones no blog de TensorFlow.
- **~ 70,000** commits.
- **> 2,200** contribuyentes.



# Comunidad de TF

2.0 is ~~coming~~  
here!

Parte de **Google ML Community** - 25 grupos [?](#)

## TensorFlow Colombia

Medellín, Colombia

1677 miembros · Grupo público [?](#)

Organizado por Yeis Taborda H. y otras 4 personas

Compartir:



# Documentación adicional

[TensorFlow](#)   [Install](#)   [Learn](#) ▾   [API](#) ▾   [Resources](#) ▾   [Community](#) ▾   [More](#) ▾    [Search](#)   [English](#) ▾   [GitHub](#)   [Sign in](#)

TensorFlow Core v2.6.0

[Overview](#)   [Python](#)   [C++](#)   [Java](#)

## Overview

[TensorFlow](#)  
[TensorFlow JavaScript](#)  
[TensorFlow Lite](#)  
[TensorFlow Extended](#)  
[Swift for TensorFlow](#)

## Versions

[Overview](#)  
2.6 (stable)  
1.15  
► More...

## Community supported languages

[Haskell](#)   
[C#](#)   
[Julia](#)

[TensorFlow](#) > [API](#) > [TensorFlow Core v2.6.0](#)

Was this helpful?

## API Documentation

TensorFlow has APIs available in several languages both for constructing and executing a TensorFlow graph. The Python API is at present the most complete and the easiest to use, but other language APIs may be easier to integrate into projects and may offer some performance advantages in graph execution.

A word of caution: the APIs in languages other than Python are not yet covered by the [API stability promises](#).

- [Python](#)
- [JavaScript](#)
- [C++](#)
- [Java](#)



# Uso de data pipelines



# En este módulo aprenderemos

- Cómo cargar bases de datos en formatos como CSV, JSON, Base64.
- Pre-procesar los datos.
- Cómo cargar dataset de Keras.
- Dataset generators.
- Cómo cargar tu propios dataset con TF.data.
- Cómo distribuir los datos.

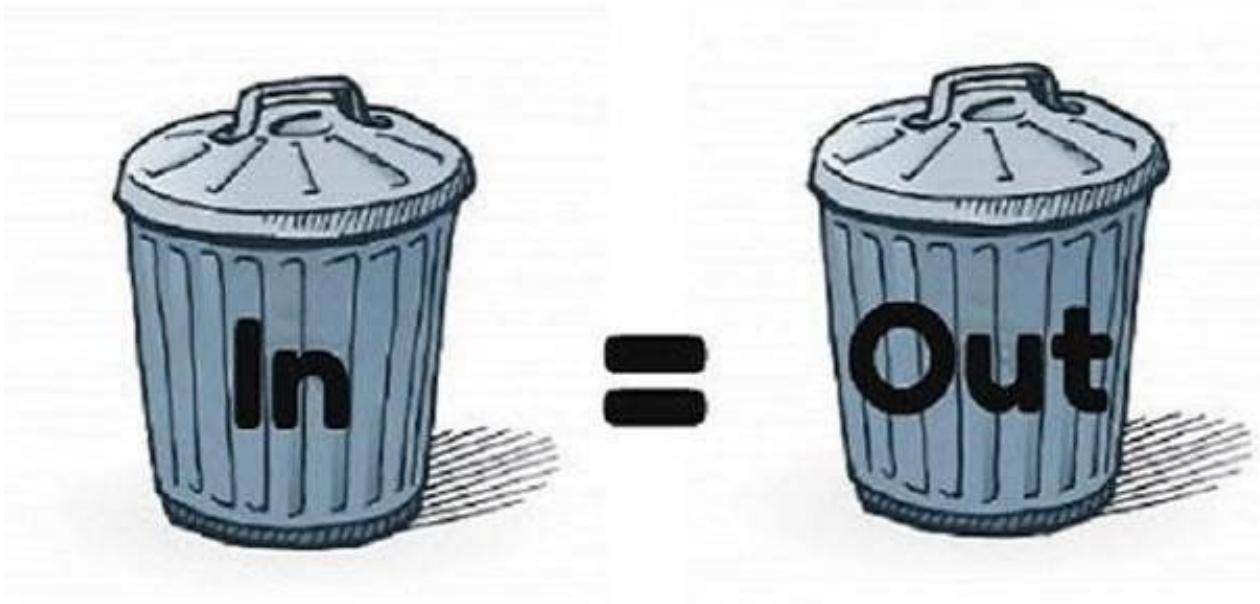


# Motor de inteligencia artificial





**Basura que entra  
= Basura que sale**



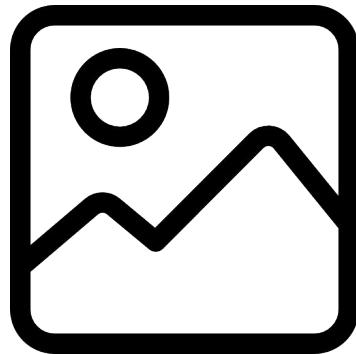


# Pre-procesar la información

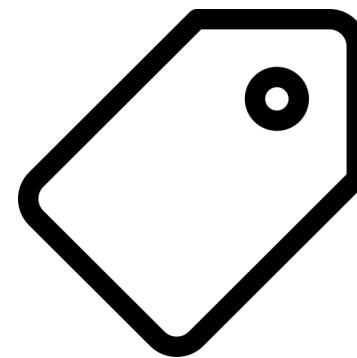




# ¿Qué compone un dataset?



Imágenes



Labels



# Cargar bases de datos JSON con URL a la nube de GCP



# Cargar bases de datos Json Base64 y .CSV



# Preprocesamiento de datos



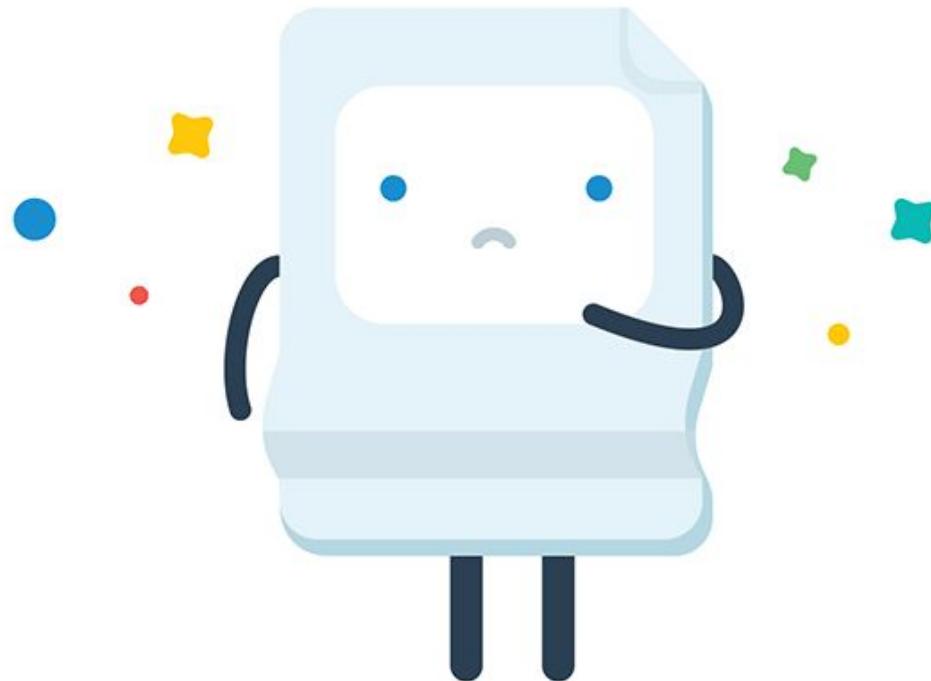
# Keras datasets



# Dataset Generators



# Memoria limitada





# Generator en Python

Es una función que devuelve un objeto, el cual te permite iterar en el mismo, sin almacenarlos a la vez en memoria.

The screenshot shows the Platzi website interface. At the top, there's a navigation bar with links for Clases, Blog, Foro, Agenda, TV, and Planes. A user profile icon shows 108 pts. On the right is a search bar. Below the navigation, there's a promotional banner for 'CONF' with the text 'Suscríbete a Expert y aprende de tecnología al mejor precio anual.' It shows a price of '\$790.000' with a Colombian flag, and a discount message 'Antes:\$929.000 Ahorras:\$139.000'. A green button says 'COMIENZA AHORA'.

[Inicio > Curso Práctico de Python: Creación de un CRUD](#)



## Yield in python

Llegó el momento en que me pregunté, ¿Qué hace `yield` en python? Es el momento de responder esa pregunta, pero para ello hay que entender qué es un **iterable** y qué es un **iterator**

[<h1>Iterable here!</h1>](#)

Se le llama así a un objeto que es capaz de retornar sus miembros uno a la vez. Entre los iterables se encuentran estructuras de secuencia como son:

- strings. str
- listas. list
- tuples. tuple



# Aprende a buscar bases de datos



# ¿Donde buscar?



kaggle

Google  
Dataset  
Search Beta





# Cómo distribuir los datos



# Distribución de bases de datos

TRAINING



VALIDATION



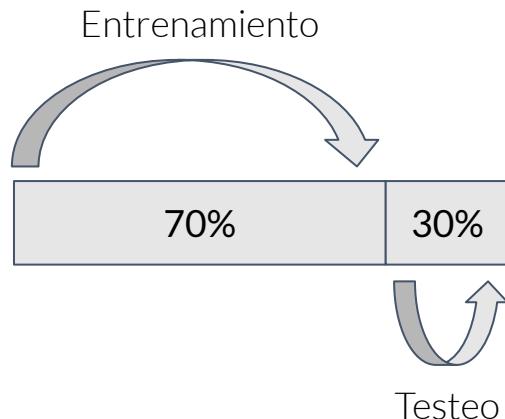
TEST



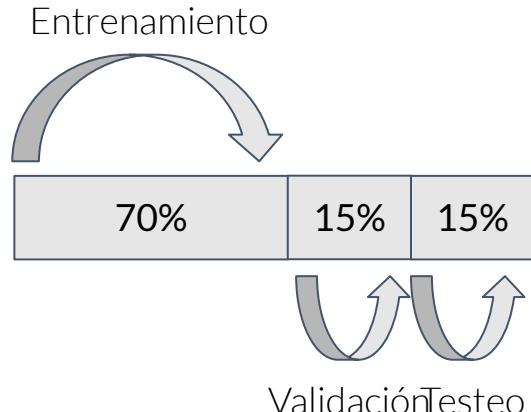


# Distribución de los datos según casos

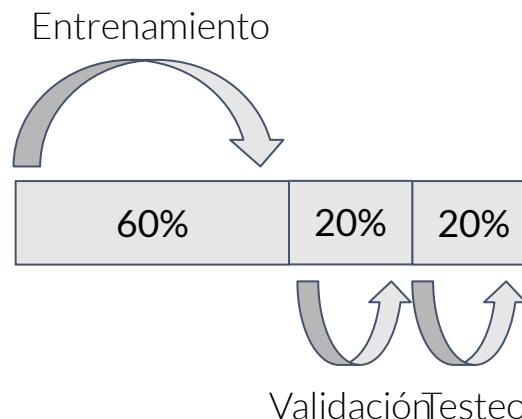
## CONFIGURACIÓN PROMEDIO



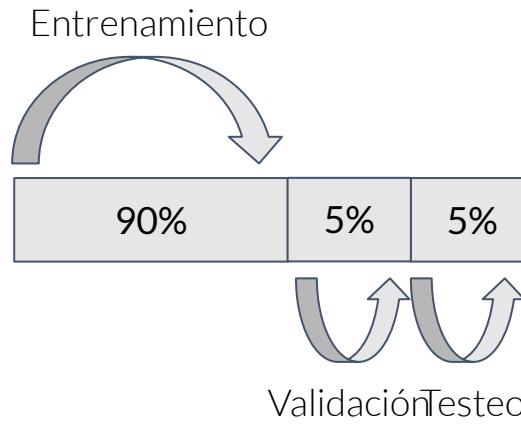
## CONFIGURACIÓN RECOMENDADA



## ANDREW NG

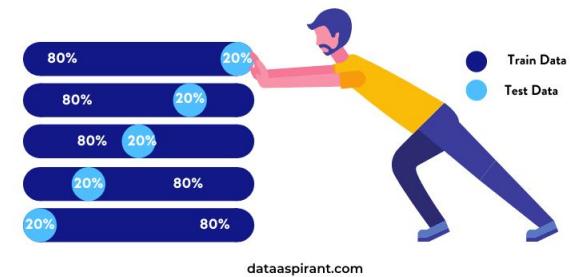


## MUCHOS DATOS



## POCOS DATOS

## Cross Validation

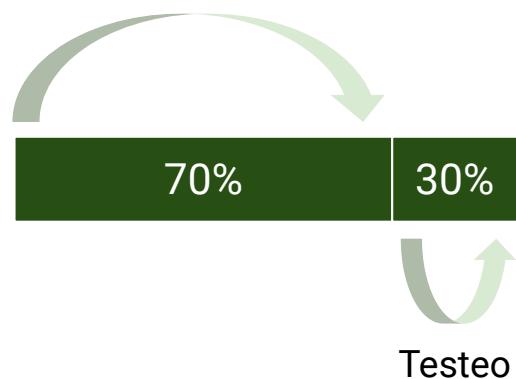




# Distribución de los datos según casos

## CONFIGURACIÓN PROMEDIO

Entrenamiento



70%

30%

70%

15%

15%

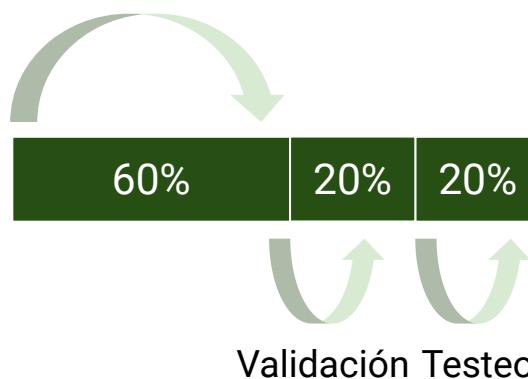
Testeo

Entrenamiento

Validación Testeo

## ANDREW NG

Entrenamiento



60%

20%

20%

Validación Testeo

## MUCHOS DATOS

Entrenamiento



90%

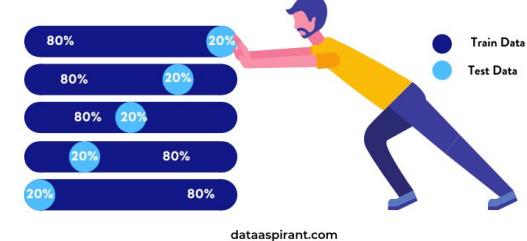
5%

5%

Validación Testeo

## POCOS DATOS

### Cross Validation



dataaspirant.com



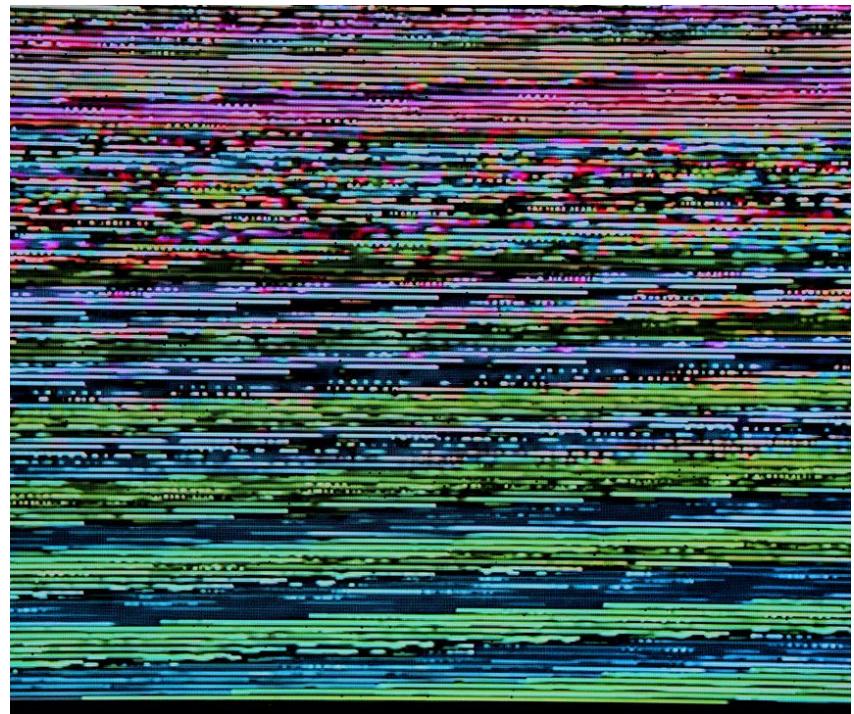
# Cuidado con confundir a tu modelo





# Errores más comunes

- Agregar datos de entrenamiento a testeo.
- Bases de datos no balanceadas en clases.
- Muy pocos datos.





Llevemos todos  
nuestros códigos a la  
acción - Crear modelo



# Optimización de precisión de modelos



# En esta sección veremos

- Underfitting y overfitting.
- Recomendaciones prácticas para ajustar mi modelo.
- Métricas para medir eficacia: callbacks.
- Monitoriza el entrenamiento de modelos con early stopping.
- Autotunning con Keras.



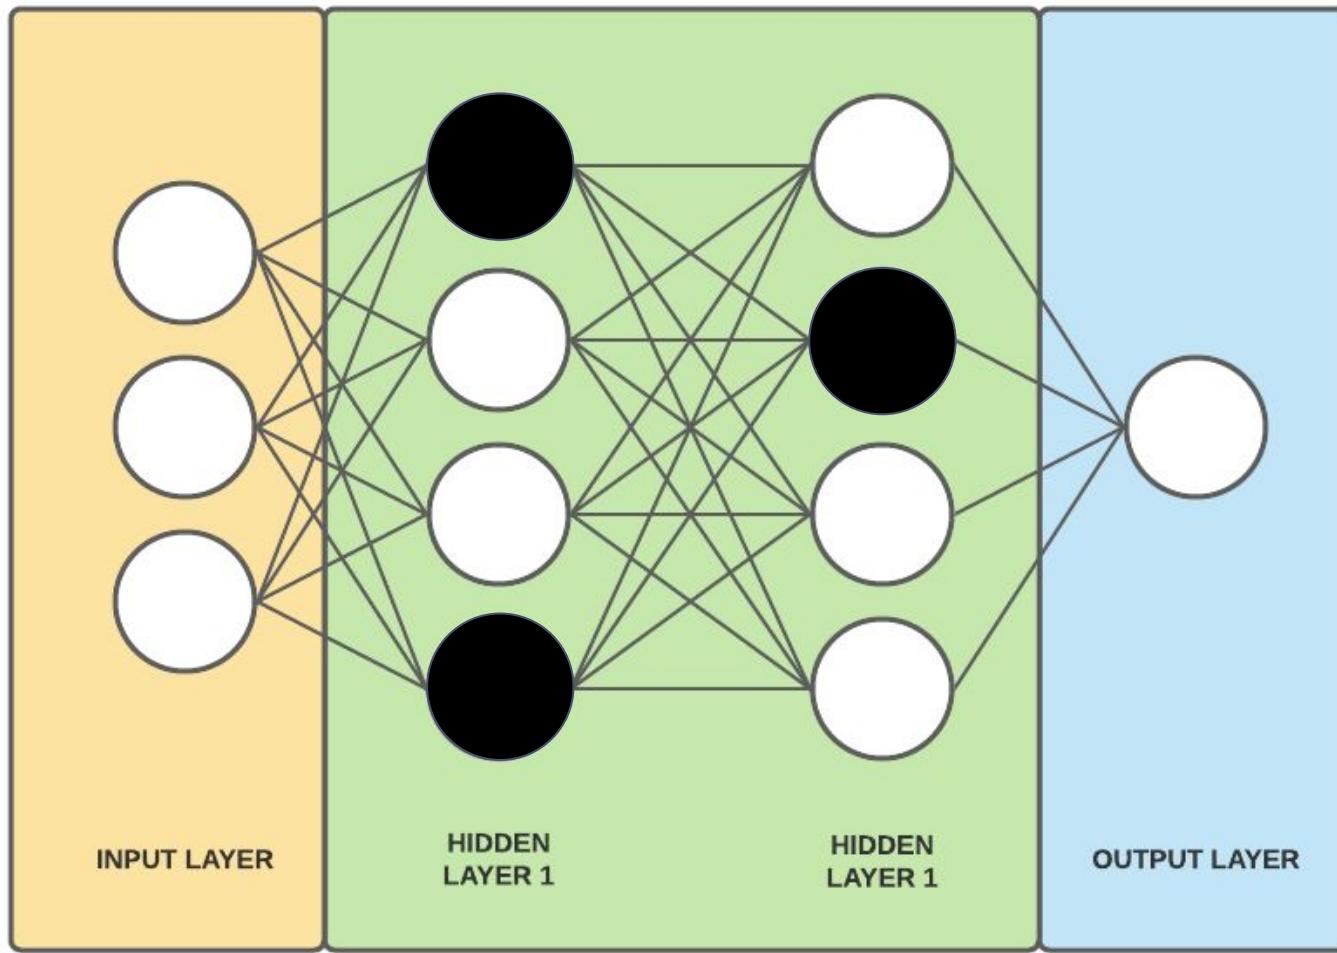


# Métodos de regularización

Overfitting y underfitting



# Dropout



*Designed by: Correlation one*



# ¿Cuál regularizador utilizar?

## Lasso L1



Datos de entrada irrelevantes.

## Ridge L2



Datos de entrada correlacionados entre ellos.

## ElasticNet L1 + L2



Gran número de atributos.

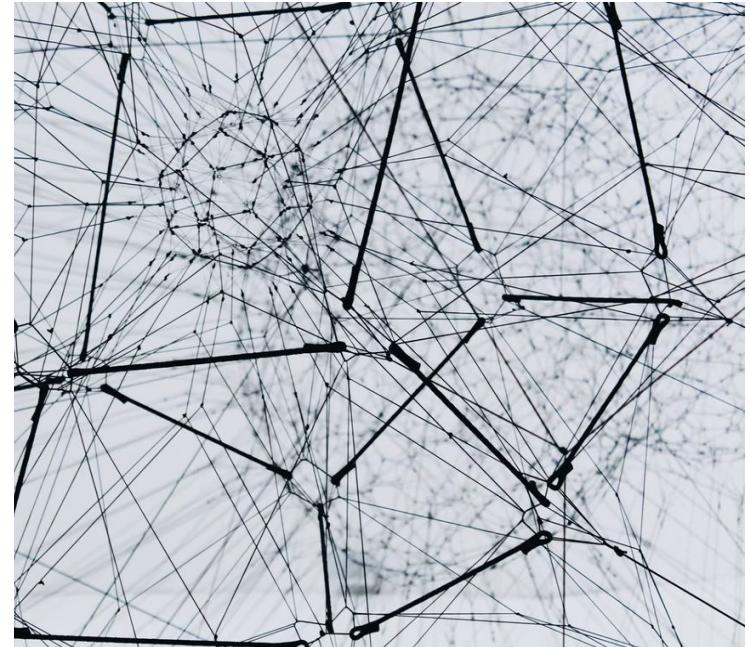


# Recomendaciones prácticas para ajustar mi modelo



# Recomendaciones para preprocesamiento

- Busca datos null.
- Archivos corruptos.
- Balancea tu base de datos.
- Aplica normalización.
- Visualiza la base de datos.
- Entiende los datos.





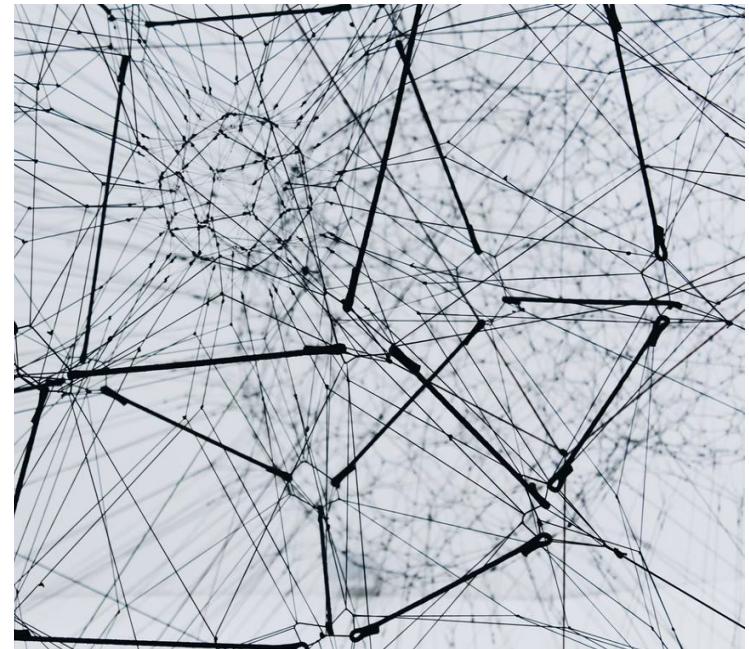
# Recomendaciones de valores

- Convoluciones (3x3)
- Pooling (2x2)
- Flatten (imágenes)
- Neuronas (64, 128, 256, 512)
- Learning Rate (0.001) ADAM
- L1\_2 (1e-5)
- Dropout (0.2)



# Recomendaciones para regularizadores

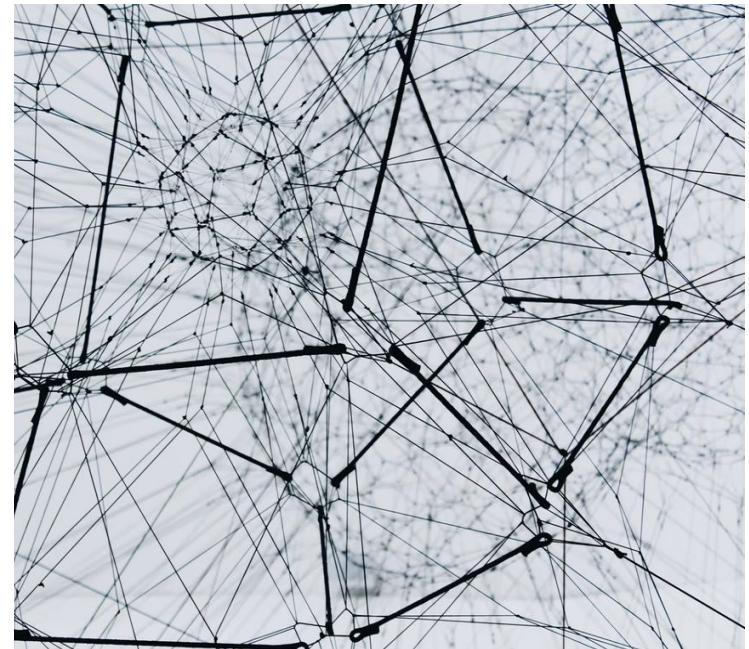
- L1, L2, L1+L2
- Agrega más datos.
- Data augmentation.
- Dropout.
- Early stopping.
- Callbacks.





# Recomendaciones funciones de activación

- Multi Clases: **Softmax**
- Binarios: **Sigmoidal**
- Regresiones: **Función lineal.**
- Clase predicciones mayor que 0: **ReLU.**





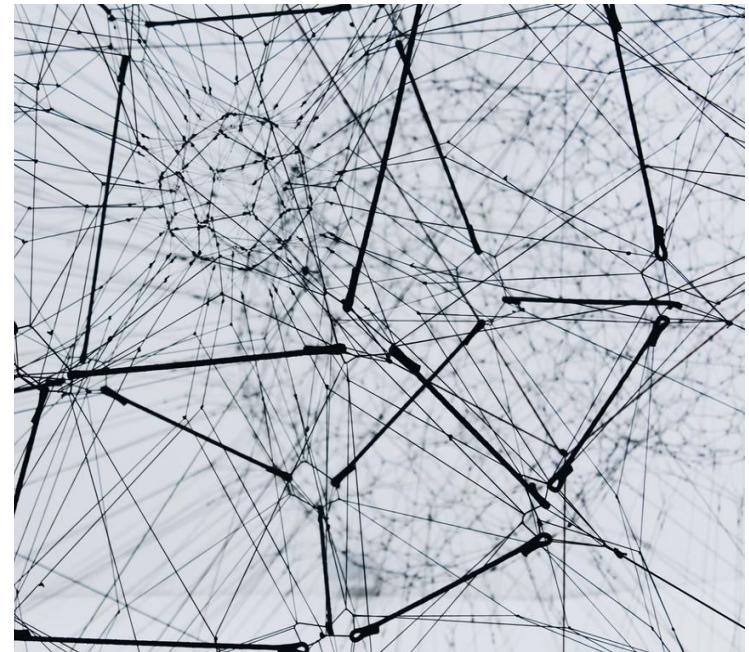
# Recomendaciones configuración de red

- Aplica capas convoluciones y poolings (1 layers).
- Tómate el tiempo de encontrar el learning rate óptimo.
- Ve almacenando en cada epoch de entrenamiento.



# Recomendaciones de la red

- Entre más capas y neuronas, no significa que es mejor la red.
- La solución no es siempre redes neuronales.
- Aliado aprendizaje por transferencia.





# Métricas para medir la eficiencia de tu modelo

Callback



# Monitoreo del entrenamiento en tiempo real

Early stopping y patience

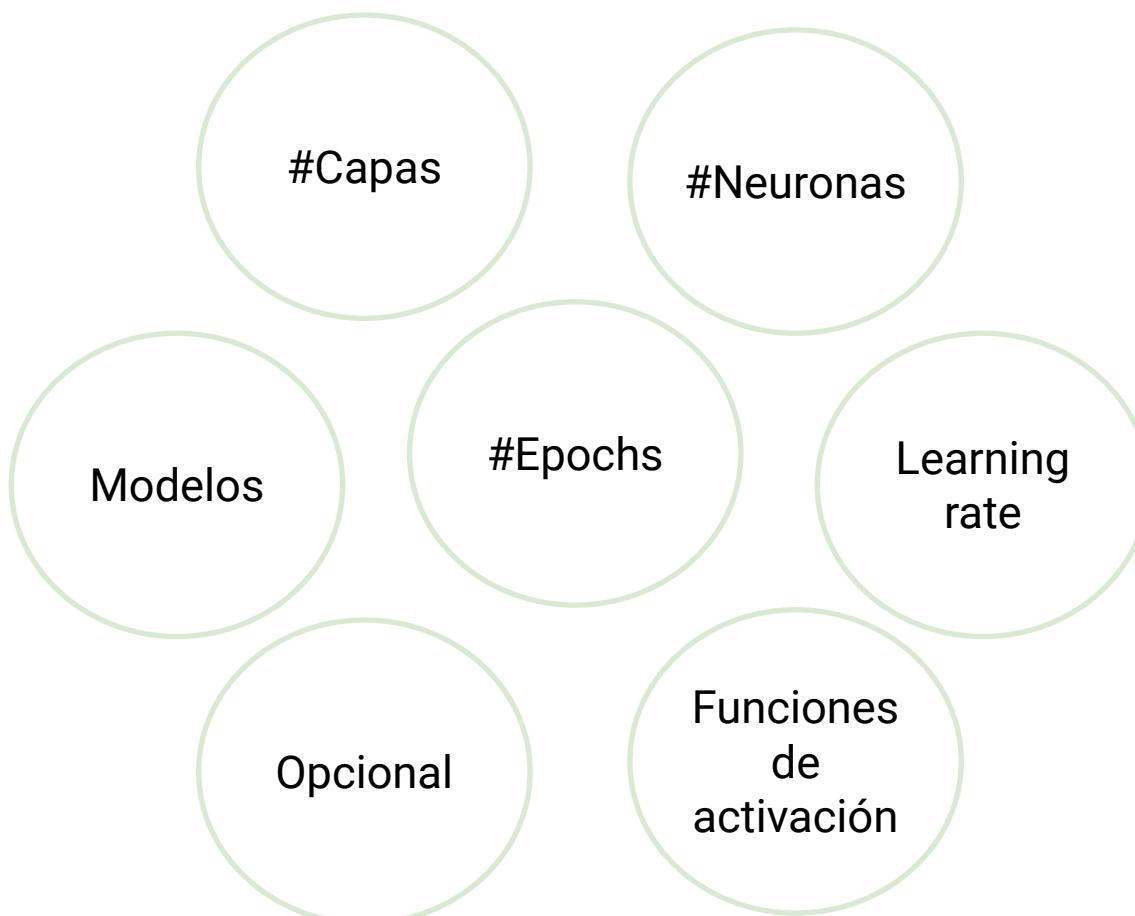


# Keras Tuner

Auto tuning



# Historia del Tunner de Keras





# Cómo se hacía antes

Results\_entrenamiento [XLSX](#)

Archivo Editar Ver Insertar Formato Datos Herramientas Ayuda

H6 | fx |

	A	B	C	D
1				
2		REDES NEURONALES		
3	Arquitectura		Loss	Accuracy
4	hp_units = hp.Int('units', min_value=32, max_value=512, step=32) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) hp_units2 = hp.Int('units', min_value=32, max_value=512, step=32) model.add(keras.layers.Dense(units=hp_units2, activation='tanh')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Flatten()) model.add(keras.layers.Dense(4))		1,636657119	0,8537979126
5	model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Flatten()) model.add(keras.layers.Dense(4))		2,440923691	0,8797893524
6	hp_units = hp.Int('units', min_value=16, max_value=512, step=16) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Flatten()) model.add(keras.layers.Dense(4))		8,43824832	0,5502405
7	hp_units = hp.Int('units', min_value=16, max_value=512, step=16) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid', kernel_regularizer=regularizers.l2(0.001))) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Dense(units=hp_units, activation='sigmoid')) model.add(keras.layers.Flatten()) model.add(keras.layers.Dense(4))		8,579723358	0,5503473282



# Keras Tuner

Buscando la mejor configuración  
para tu modelo



# Almacenamiento y carga de modelos

Pesos y arquitectura



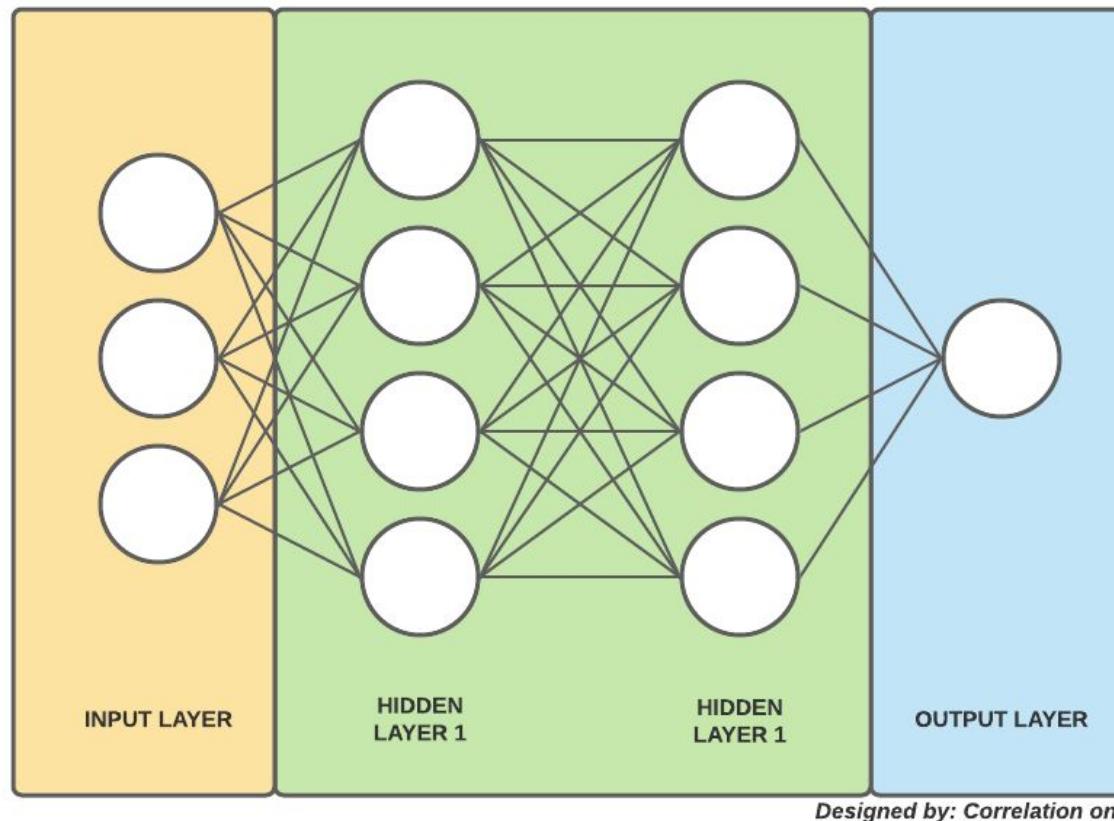
# En esta sección aprenderemos

- Almacenamiento y carga de modelos: pesos y arquitectura.
- Criterios para almacenar los modelos.





# Componentes relevantes de un modelo



Arquitectura

Pesos

Etiquetas (Labels)  
Aprendizaje por transferencia.



# Criterios para almacenar los modelos



# Introducción al aprendizaje por transferencia



# En este módulo aprenderemos

- Introducción al aprendizaje por transferencia.
- Cómo cargar sistemas pre-entrenados en Keras.
- Cómo utilizar sistemas pre-entrenados desde los repositorios de TensorFlow Hub.
- Aplicar modelo pre-entrenado a nuestro proyecto.



# Transferir conocimiento





# Transferir características principales



Modelo Matemático



Modelo Matemático





# Bases de datos de entrenamiento



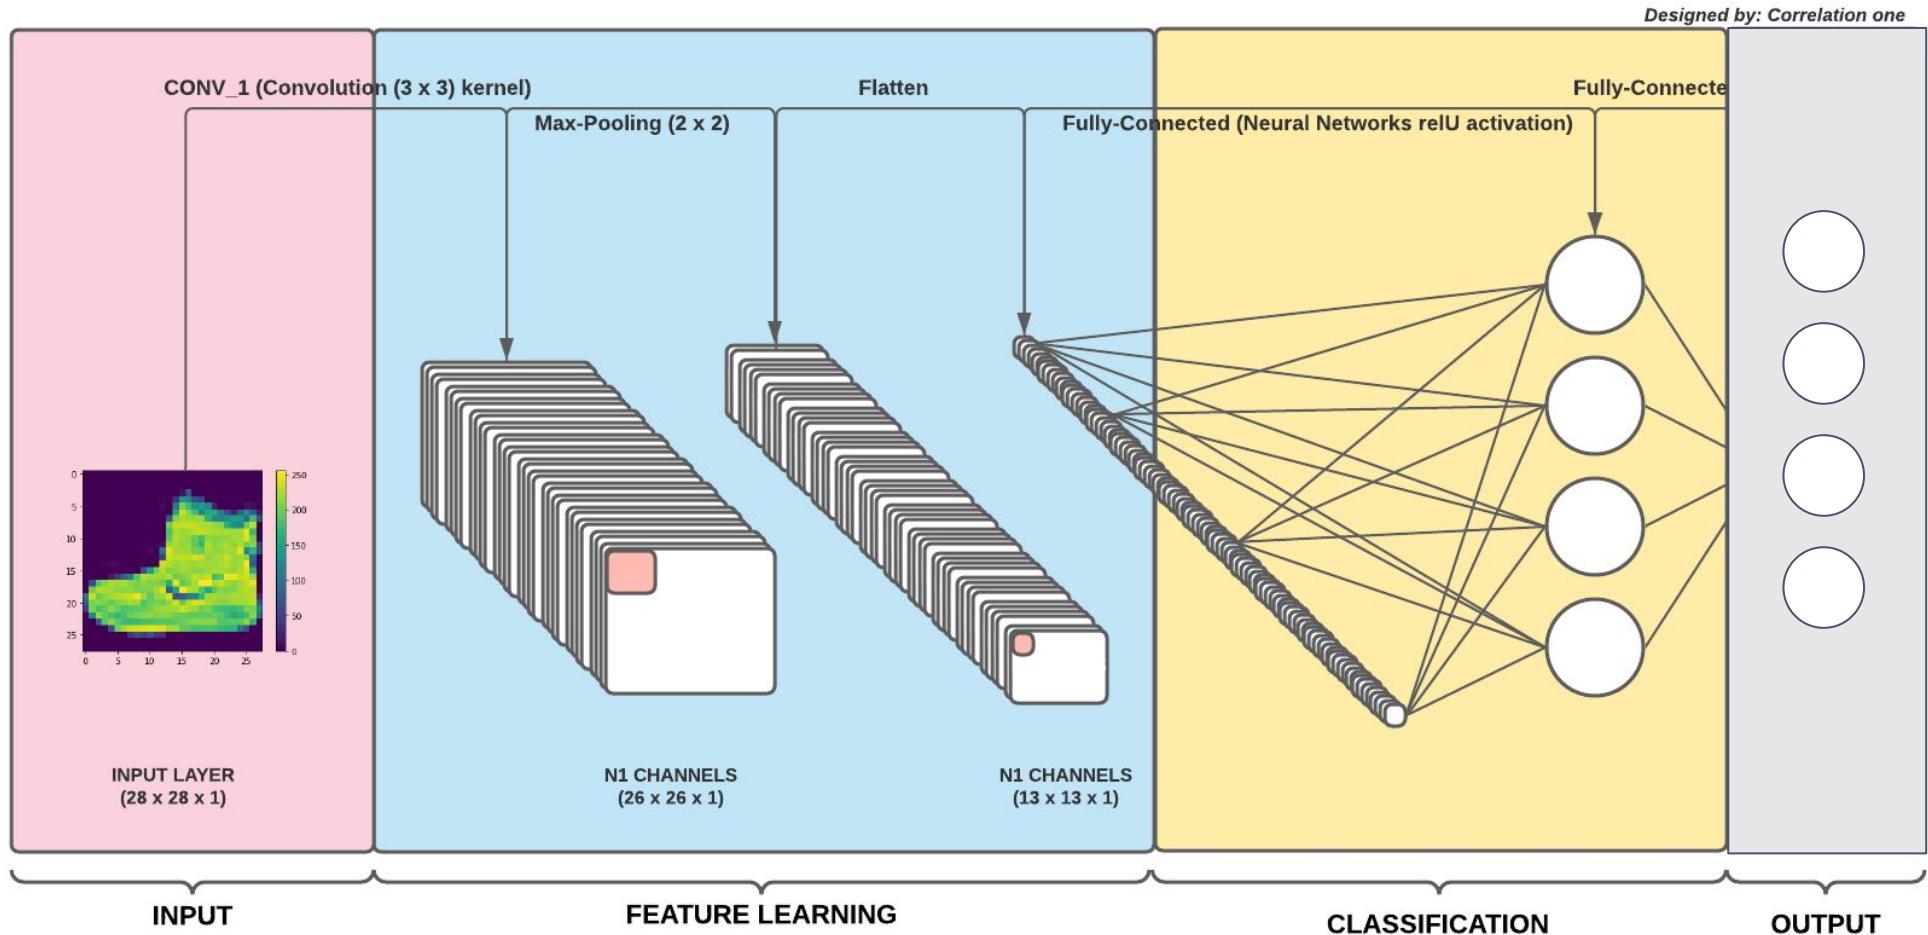
## OpenAI

GPT-3, an autoregressive language model with 175 billion parameters



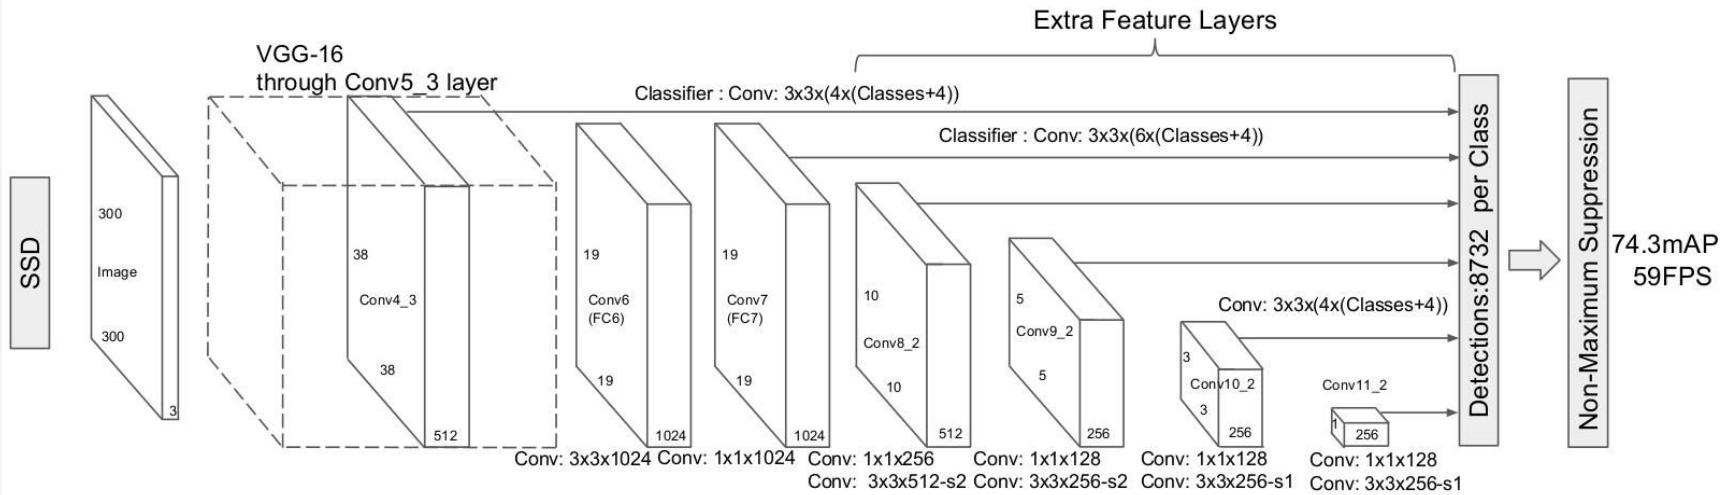


# ¿Cómo funcionan?





# Ejemplo de infraestructuras de una red convolucional



Fuente: <https://arxiv.org/abs/1512.02325>



```
from tensorflow.keras.layers import Dense
from tensorflow.keras.Model import Sequential
from tensorflow.keras.applications.inception_v3 import InceptionV3

## Url inception model.
weights_file = '/tmp/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5'

new_model = Sequential()
new_model.add(InceptionV3(include_top = False,
                           weights = weights_file))

new_model.add(Dense(num_classes, activation = 'softmax'))
new_model.layers[0].trainable = False
```



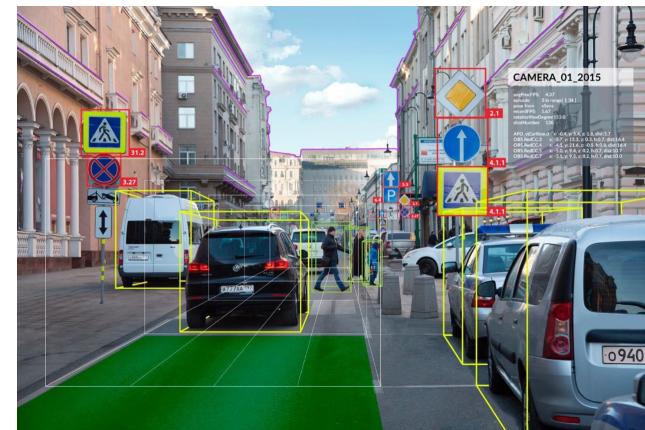
# Cuándo utilizar aprendizaje por transferencia



# Cuándo utilizar modelos pre-entrenados



Procesamiento de  
lenguaje natural

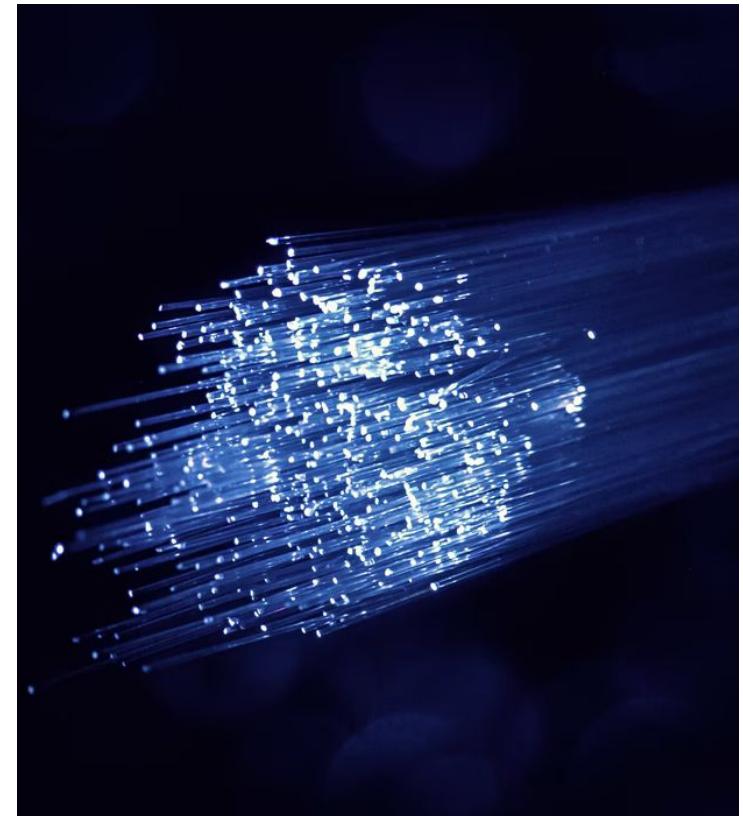


Visión  
computarizada



# ¿Por qué utilizarlos?

- Se necesita poca data.
- Te permite generar iteraciones muy rápidas.
- Generaliza muy bien los modelos.
- Transferir extracción de características principales.





# Recomendaciones de artículos pre-entrenados

## YOLOv3: An Incremental Improvement

Joseph Redmon, Ali Farhadi  
University of Washington

### Abstract

We present some updates to YOLO. We made a bunch of little design changes to make it better. We also trained this new network that's pretty swell. It's a little bigger than last time but more accurate. It's still fast though, don't worry. At 32x32, it runs at 30ms on a Titan X. And as always, as far as I know, it has three times faster. When we look at the old 5. IOU@P detection metric, YOLOv3 is quite good. It achieves 57.9 AP<sub>50</sub> in 51 ms on a Titan X, compared to 57.5 AP<sub>50</sub> in 198 ms by RetinaNet. Similar performance but 3.8x faster. As always, all the code is online at <https://pjreddie.com/yolo/>.

### 1. Introduction

Sometimes you just kinda phone it in for a year, you know? I didn't do a whole lot of research this year. Spent a lot of time on Twitter. Played around with GANs a little. I had a little time left over from last year [1][2][3]. I managed to make some improvements, though. Most notably, nothing like super interesting, just a bunch of small changes that make it better. It also helped out with other people's research a little.

Actually, I'm not taking us here today. We have known! I didn't do a whole lot of research this year. Spent a lot of time on Twitter. Played around with GANs a little. I had a little time left over from last year [1][2][3]. I managed to make some improvements, though. Most notably, nothing like super interesting, just a bunch of small changes that make it better. It also helped out with other people's research a little.

Actually, I'm not taking us here today. We have

known! I didn't do a whole lot of research this year. Spent a lot of time on Twitter. Played around with GANs a little. I had a little time left over from last year [1][2][3]. I managed to make some improvements, though. Most notably, nothing like super interesting, just a bunch of small changes that make it better. It also helped out with other people's research a little.

Following YOLO9000 our system predicts bounding boxes using dimension clusters as anchor boxes [13]. The network predicts 4 coordinates for each bounding box,  $t_x$ ,  $t_y$ ,  $t_w$ ,  $t_h$ . If the cell is offset from the top-left corner of the image by  $(c_x, c_y)$  and the bounding box prior has width and height  $p_w$ ,  $p_h$ , then the predictions correspond to:

$$\begin{aligned} b_x &= \sigma(t_x) + c_x \\ b_y &= \sigma(t_y) + c_y \\ b_w &= p_w e^{t_w} \\ b_h &= p_h e^{t_h} \end{aligned}$$

During training we use sum of squared error loss. If the ground truth for some coordinate prediction is  $t_a$ , our gradient is the ground truth value (computed from the ground truth box) minus our prediction:  $t_a - t_a$ . This ground truth value can be easily computed by inverting the equations above.

YOLOv3 predicts an objectness score for each bounding

## ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky  
University of Toronto  
kriz@cs.toronto.ca

Ilya Sutskever  
University of Toronto  
ilya@cs.toronto.ca

Geoffrey E. Hinton  
University of Toronto  
hinton@cs.toronto.ca

### Abstract

We trained a large, deep convolutional neural network to classify the 1.2 million high-resolution images in the ImageNet LSVRC-2010 dataset using 1000 different classes. On this test set, we show that top-5 error rates of 15.3% and 17.0% which is considerably better than the previous state-of-the-art. The neural network, which has 60 million parameters and 650,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and three fully connected layers. A remarkable aspect of the network is its speed: it can classify a single image in about 15 minutes on a single GPU. To achieve this speed, we used non-saturating neurons and a very efficient GPU implementation of the convolution operation. To reduce overfitting in the fully-connected layers we employed a recently-developed regularization method called "dropout" that proved to be very effective. We also entered a variant of this model in the ILSVRC-2012 competition and achieved a winning top-5 test error rate of 15.3%, compared to 26.2% achieved by the second-best entry.

### 1 Introduction

Current approaches to object recognition make essential use of machine learning methods. To improve their performance, we can collect larger datasets, learn more powerful models, and use better techniques for preventing overfitting. Until recently, datasets of labeled images were relatively small — on the order of tens of thousands of images (e.g., NORB [16], Caltech-10/256 [8, 9], and CIFAR-10/100 [12]). Simple recognition tasks can be solved quite well with datasets of this size, especially if they are combined with well-designed training data. For example, the current best error rate on the MNIST digit recognition task (0.3%) approaches the theoretical limit [4]. But objects in realistic settings exhibit considerable variability, so to learn to recognize them it is necessary to use much larger training sets. And indeed, the shortcomings of small image datasets have already been noted (e.g., Poggio et al. [14]), but this has only recently become feasible to collect labeled datasets with millions of images. The two largest datasets include LabelMe [23], which consists of hundreds of thousands of fully-segmented images, and ImageNet [6], which consists of over 15 million labeled high-resolution images in over 22,000 categories.

To learn about thousands of objects from millions of images, we need a model with a large learning capacity. However, the immense complexity of the object recognition task means that the problem cannot be specified even by a dataset as large as ImageNet, so our model should also have lots of prior knowledge to compensate for all the data we don't have. Convolutional neural networks

# YoloV3

# AlexNet



# ¿Cómo saber cuál modelo seleccionar?



Accuracy (precisión)

Coste computacional

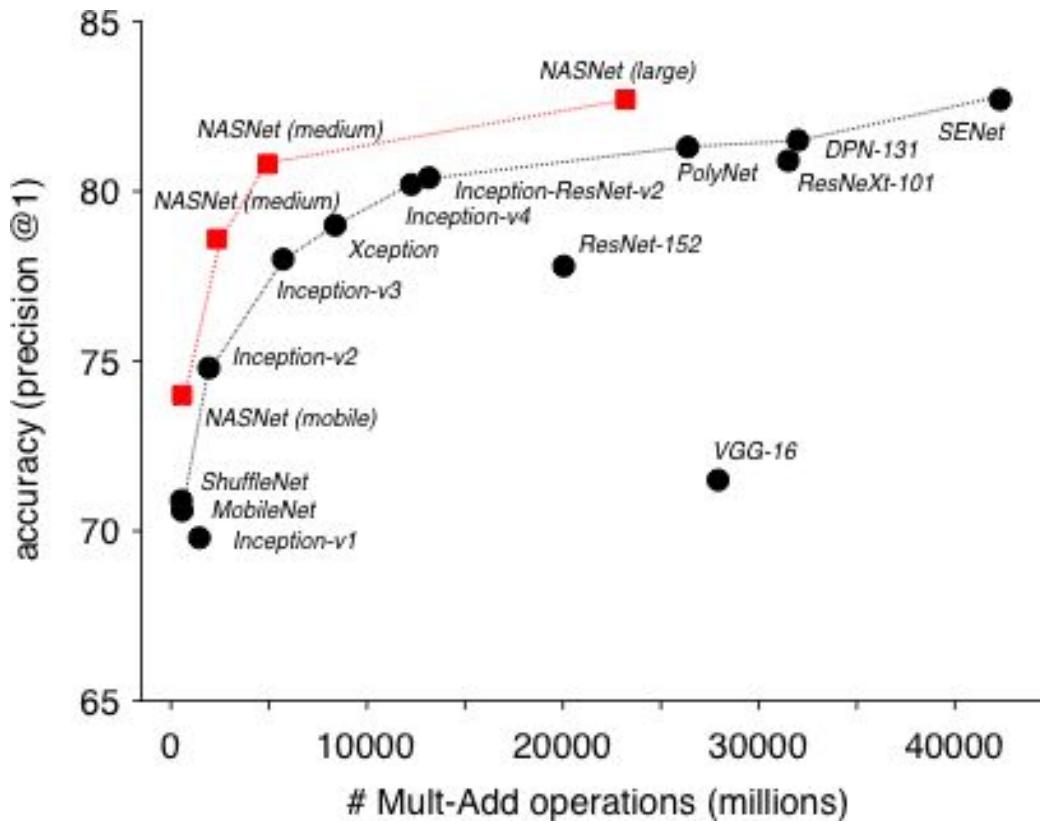
```
CPU[|||] 1.8% Tasks: 23; 1 running
Mem[||||| 41.1M/489M] Load average: 0.43 0.33 0.13
Swap[OK/OK] Uptime: 00:02:22

PID USER NI VIRT RES SHR S CPU_MEM% TIME+ Command
 1 root 20 0 37636 5644 3936 S 0.0 1.1 0:03.64 /sbin/init
1044 root 20 0 35272 2460 8124 S 0.0 1.7 0:00.48 /lib/systemd/journald
1087 root 20 0 94772 1556 1380 S 0.0 0.3 0:00.00 /bin/lvmetad -f
1099 root 20 0 42768 3928 2772 S 0.0 0.8 0:00.27 /lib/systemd/systemd-udevd
1255 systemd-t 20 0 98M 2484 2248 S 0.0 0.5 0:00.01 /lib/systemd/systemd-timesyncd
1394 daemon 20 0 26044 2084 1884 S 0.0 0.4 0:00.00 /usr/sbin/atac -f
1400 root 20 0 181M 10404 9276 S 0.0 2.1 0:00.02 /usr/lib/dbus-daemon --system --address=systemd: --nofor
1418 root 20 0 28549 2840 2536 S 0.0 0.6 0:00.03 /lib/systemd/systemd-logind
1421 root 20 0 29296 2864 2592 S 0.0 0.6 0:00.00 /usr/sbin/cron -f
1423 syslog 20 0 250M 3140 2548 S 0.0 0.6 0:00.10 /usr/sbin/rsyslogd -n
1427 root 20 0 4400 1372 1280 S 0.0 0.3 0:00.00 /usr/sbin/acpid
1429 root 20 0 95360 1392 1256 S 0.0 0.3 0:00.00 /usr/bin/lxcfs /var/lib/lxcfs/
1447 root 20 0 269M 5992 5352 S 0.0 1.2 0:00.05 /usr/lib/accountservice/accounts-daemon
1475 root 20 0 13376 164 20 5 S 0.0 0.0 0:00.00 /sbin/mdadm --monitor --pid-file /run/mdadm/monitor.pid
1476 root 20 0 270M 5868 5180 S 0.0 1.2 0:00.00 /usr/lib/polkitkit-1/polkitd --no-debug
2340 root 20 0 656M 5972 5268 S 0.0 1.1 0:00.08 /usr/sbin/sshd -D
2602 root 20 0 152 156 36 S 0.0 0.0 0:00.00 /sbin/iscsicd
2363 root 10 0 5724 5524 2432 S 0.0 0.7 0:00.05 /sbin/iscsicd
2420 root 20 0 16228 1568 1424 S 0.0 0.3 0:00.01 /sbin/agetty --noclear tty1 linux
2567 root 20 0 95460 6828 5880 S 0.0 1.4 0:00.12 sshd: root@pts/0
2602 root 20 0 22860 4476 2604 S 0.0 0.9 0:00.22 -bash
2857 root 20 0 26480 4224 3236 R 0.9 0.8 0:00.16 htop

F1Help F2Setup F3Search F4Filter F5Spec F6SortByF7Nice F8Nice F9Kill F10Quit
```



# ¿Cómo seleccionar cual modelo pre-entrenado usar?



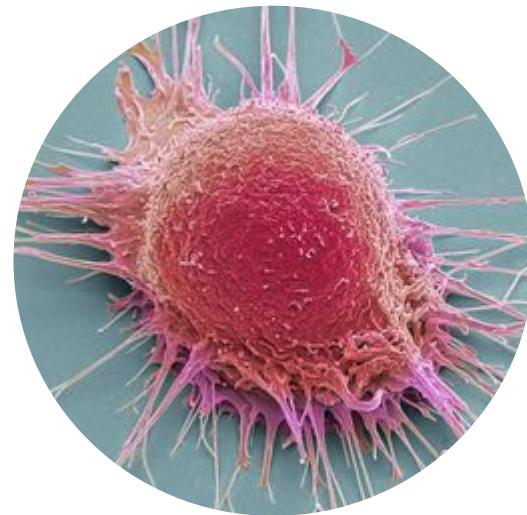
Fuente: <https://ai.googleblog.com/2017/11/automl-for-large-scale-image.html>



# Reto



Dron detectando y  
persiguiendo aves.



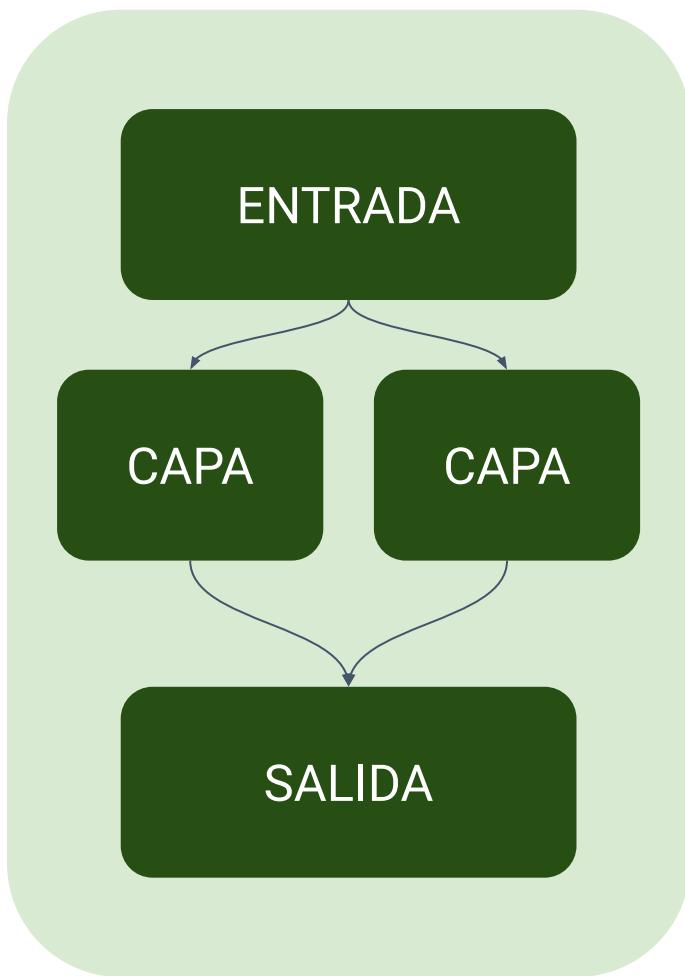
Detección de  
cáncer.



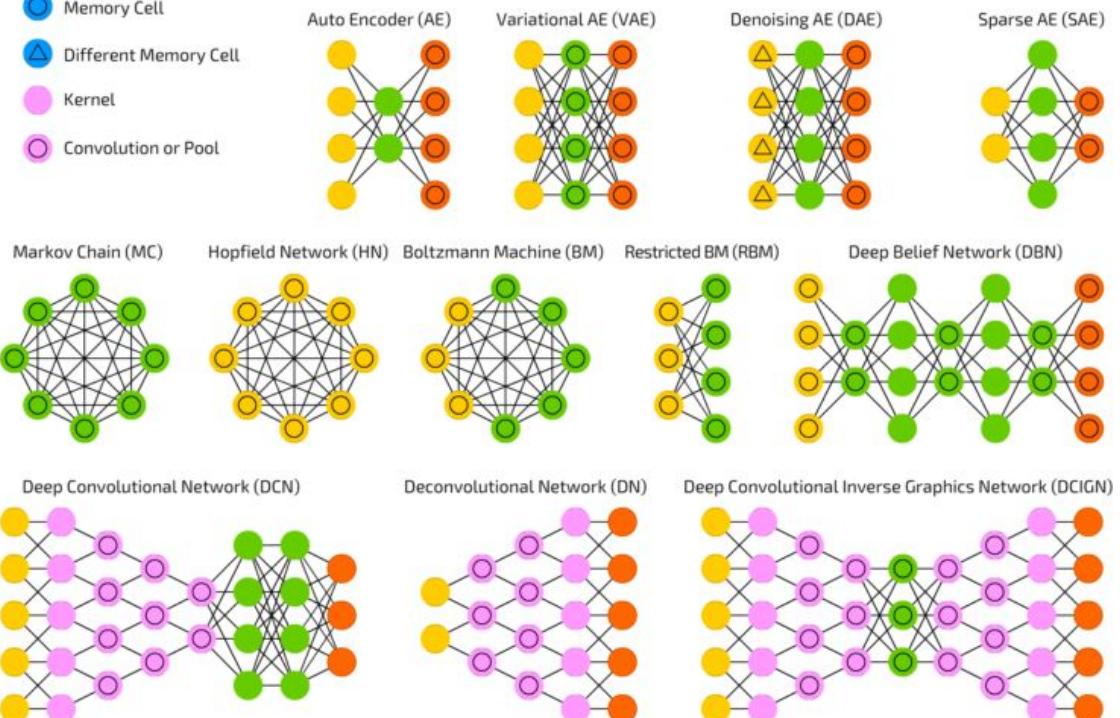
# Carga de sistemas pre-entrenados en Keras



# API Funcional Keras



- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

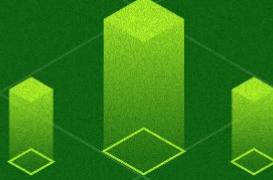




# API funcional de Keras



# Uso sistemas pre-entrenados de TensorFlow Hub



# Introducción a variables relevantes del TensorBoard



# En esta sección aprenderemos

- Introducción a variables relevantes del TensorBoard.
- Análisis y publicación de resultados del entrenamiento.
- Introducción para poner modelos en producción.
- Cómo poner en producción.



# Introducción a TensorBoard

- Herramienta para visualizar resultados.
- Entender el flujo y los tensores.
- Debuggear tu modelo.





# Ejemplo práctico de TensorBoard

TensorBoard      SCALARS      GRAPHS      DISTRIBUTIONS      HISTOGRAMS      PROJECTOR      INACTIVE      C      G      ?

Show data download links

Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing: 0,6

Horizontal Axis: STEP (selected), RELATIVE, WALL

Runs: Write a regex to filter runs

dnn

average\_loss

average\_loss

run to download CSV JSON

dnn/dnn/hiddenlayer\_0/fraction\_of\_zero\_values

dnn/dnn/hiddenlayer\_1/fraction\_of\_zero\_values

dnn/dnn/logits/fraction\_of\_zero\_values

1 3



# Análisis y publicación de resultados del entrenamiento



# Introducción al despliegue de modelos en producción



# Production





# Ejemplos de producción

Nube



Google Cloud



Sistemas embebidos





# Ejemplos de producción

Local



USB Accelerator



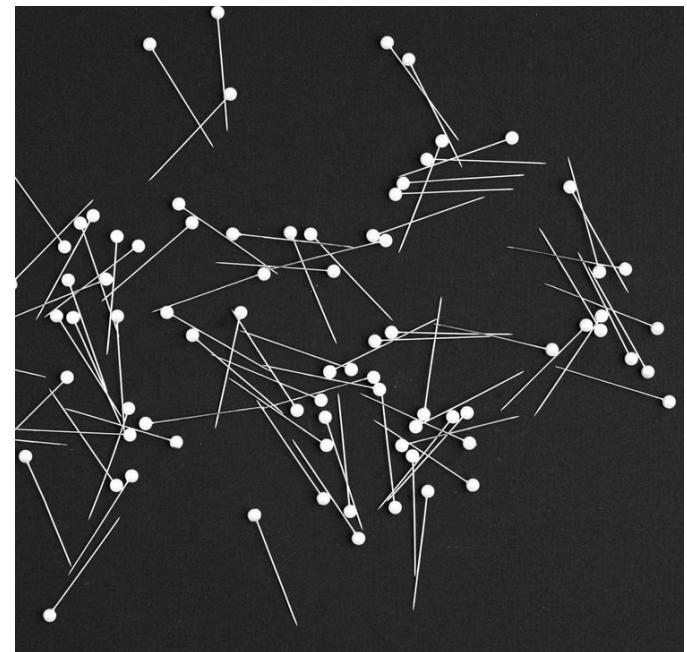


# Siguientes pasos con deep learning



# Qué más aprender

- Redes neuronales convolucionales.
- Data Augmentation.
- Formato TF Records para entrenar modelos.
- Despliegue de modelos a producción.
- Visión computarizada.



A portrait of a young man with short dark hair and round glasses, smiling broadly. He is wearing a long-sleeved, button-up shirt in a muted olive-green color. The background is plain white.

[www.switchai.co](http://www.switchai.co)